

BlackJack

December 15, 2018

1 Monte Carlo Methods for Blackjack

In this notebook I wrote my own implementations of many Monte Carlo algorithms, in order to teach an agent how to play Blackjack.

1.0.1 Explore BlackjackEnv

This is an instance of the [Blackjack](#) environment.

```
In [1]: import matplotlib
import numpy as np
import sys
import gym

%matplotlib inline

env = gym.make('Blackjack-v0')

print(env.observation_space)
print(env.action_space)
```

```
Tuple(Discrete(32), Discrete(11), Discrete(2))
Discrete(2)
```

Each state is a 3-tuple of: - the player's current sum $\in \{0, 1, \dots, 31\}$, - the dealer's face up card $\in \{1, \dots, 10\}$, and - whether or not the player has a usable ace (no = 0, yes = 1).

The agent has two potential actions:

- STICK = 0
- HIT = 1

1.0.2 Play Blackjack with a random policy.

This is designed to visualize the output that is returned as the agent interacts with the environment.

```
In [2]: for i_episode in range(3):
        state = env.reset()
        while True:
            print(f"State: {state}")
            action = env.action_space.sample()
            print(f"Action: {action}")
            state, reward, done, info = env.step(action)
            if action:
                print(f"New state: {state}")
            if done:
                print('End game! Reward: ', reward)
                print('You won!!!\n') if reward > 0 else print('You lost...\n')
                break
```

```
State: (15, 4, True)
Action: 0
End game! Reward: -1.0
You lost...
```

```
State: (11, 10, False)
Action: 1
New state: (13, 10, False)
State: (13, 10, False)
Action: 1
New state: (20, 10, False)
State: (20, 10, False)
Action: 0
End game! Reward: 0.0
You lost...
```

```
State: (14, 10, False)
Action: 1
New state: (24, 10, False)
End game! Reward: -1
You lost...
```

1.0.3 Generate an episode from action-value function estimate following the epsilon-greedy policy

The function accepts as **input**: - env: This is an instance of OpenAI Gym's Blackjack environment.
 - Q: This is a dictionary (of one-dimensional arrays) where $Q[s][a]$ is the estimated action value corresponding to state s and action a . - epsilon: Value between 0 and 1, (probability of choosing a random action instead the Greedy action). - nA: Number of actions the agent can choose.

It returns as **output**: - episode: This is a list of (state, action, reward) tuples (of tuples) and corresponds to $(S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T)$, where T is the final time step. In particular, `episode[i]` returns (S_i, A_i, R_{i+1}) , and `episode[i][0]`, `episode[i][1]`, and `episode[i][2]`

return S_i , A_i , and R_{i+1} , respectively.

```
In [3]: def generate_episode_from_Q(env, Q, epsilon, nA):
        episode = []
        state = env.reset()
        while True:
            action = np.random.choice(np.arange(nA), p=get_probs(Q[state], epsilon, nA)) \
                                if state in Q else env.action_space.sample()
            next_state, reward, done, info = env.step(action)
            episode.append((state, action, reward))
            state = next_state
            if done:
                break
        return episode
```

1.0.4 Calculate the action probabilities corresponding to epsilon-greedy policy

The function accepts as **input**: - Q_s : This is a dictionary (of one-dimensional arrays) where $Q[s][a]$ is the estimated action value corresponding to state s and action a . - ϵ : Value between 0 and 1, (probability of choosing a random action instead the Greedy action). - nA : Number of actions the agent can choose.

It returns as **output**: - $policy_s$: This is a list containing the probability of taking each action.

```
In [4]: def get_probs(Q_s, epsilon, nA):
        policy_s = np.ones(nA) * epsilon / nA
        best_a = np.argmax(Q_s)
        policy_s[best_a] = 1 - epsilon + (epsilon / nA)
        return policy_s
```

1.0.5 Update the action-value function estimate using the most recent episode

This algorithm has four arguments: - env : This is an instance of an OpenAI Gym environment. - $episode$: This is the number of episode that is generated through agent-environment interaction. - Q : This is a dictionary (of one-dimensional arrays) where $Q[s][a]$ is the estimated action value corresponding to state s and action a . - α : This is the step-size parameter for the update step. - γ : This is the discount rate. It must be a value between 0 and 1 inclusive.

The algorithm returns as output: - Q : This is the updated Q .

```
In [5]: def update_Q_alpha(env, episode, Q, alpha, gamma):
        states, actions, rewards = zip(*episode)
        # prepare for discounting
        discounts = np.array([gamma**i for i in range(len(rewards)+1)])
        for i, state in enumerate(states):
            # update Q using MC constant alpha
            # FORMULA ->  $Q(S_t|A_t) = Q(S_t|A_t) + \alpha * (G_t - Q(S_t|A_t))$ 
            Q[state][actions[i]] = Q[state][actions[i]] + alpha*(sum(rewards[i:]*discounts
        return Q
```

1.0.6 MC Control: Constant- α

This algorithm has four arguments: - env: This is an instance of an OpenAI Gym environment. - num_episodes: This is the number of episodes that are generated through agent-environment interaction. - alpha: This is the step-size parameter for the update step. - gamma: This is the discount rate. It must be a value between 0 and 1 inclusive.

The algorithm returns as output: - Q: This is a dictionary (of one-dimensional arrays) where $Q[s][a]$ is the estimated action value corresponding to state s and action a . - policy: This is a dictionary where $policy[s]$ returns the action that the agent chooses after observing state s .

```
In [6]: from collections import defaultdict
```

```
def mc_control_alpha(env, num_episodes, alpha, gamma=1.0):
    nA = env.action_space.n
    # fraction of num_episodes in order to have 0.1 <= epsilon < 1
    fraction = num_episodes // 10
    # initialize empty dictionary of arrays
    Q = defaultdict(lambda: np.zeros(nA))
    # loop over episodes
    for i_episode in range(1, num_episodes+1):
        # monitor progress
        if i_episode % 10000 == 0:
            print("\rEpisode {}/{}.".format(i_episode, num_episodes), end="")
            sys.stdout.flush()
        # set the value of epsilon, after 1kk episodes it remains static to 0.1
        epsilon = 1.0/((i_episode/fraction)+1)
        if i_episode > 1000000:
            epsilon = 0.1
        # generate an episode by following epsilon-greedy policy
        episode = generate_episode_from_Q(env, Q, epsilon, nA)
        # update the action-value function estimate using the episode
        Q = update_Q_alpha(env, episode, Q, alpha, gamma)
        # determine the policy corresponding to the final action-value function estimate
        policy = dict((k,np.argmax(v)) for k, v in Q.items())
    return policy, Q
```

Obtain the estimated optimal policy and action-value function.

```
In [7]: # obtain the estimated optimal policy and action-value function
        policy_alpha, Q_alpha = mc_control_alpha(env, 2500000, 0.005, gamma=0.85)
```

Episode 2500000/2500000.

1.0.7 Plot the corresponding state-value function.



```
In [8]: from plot_utils import plot_blackjack_values
```

```

# obtain the state-value function
V_alpha = dict((k,np.max(v)) for k, v in Q_alpha.items())

# plot the state-value function
plot_blackjack_values(V_alpha)

```



Visualize the policy that is estimated to be optimal.

```



In [9]: from plot_utils import plot_policy

```


```

# plot the policy
plot_policy(policy_alpha)

```



True Optimal Policy



In the following picture there is the **true** optimal policy, is not exactly like the optimal policy found by the agent but it is very similar.

In []: